

In [14]:

```
import numpy as np
from matplotlib import pyplot as plt
%matplotlib inline
```

## The MNIST dataset

The MNIST dataset is composed of 70,000  $28 \times 28$  grayscale images of handwritten digits. It is represented as a  $70000 \times 28 \times 28$  numpy array (a "3d matrix").

In [15]:

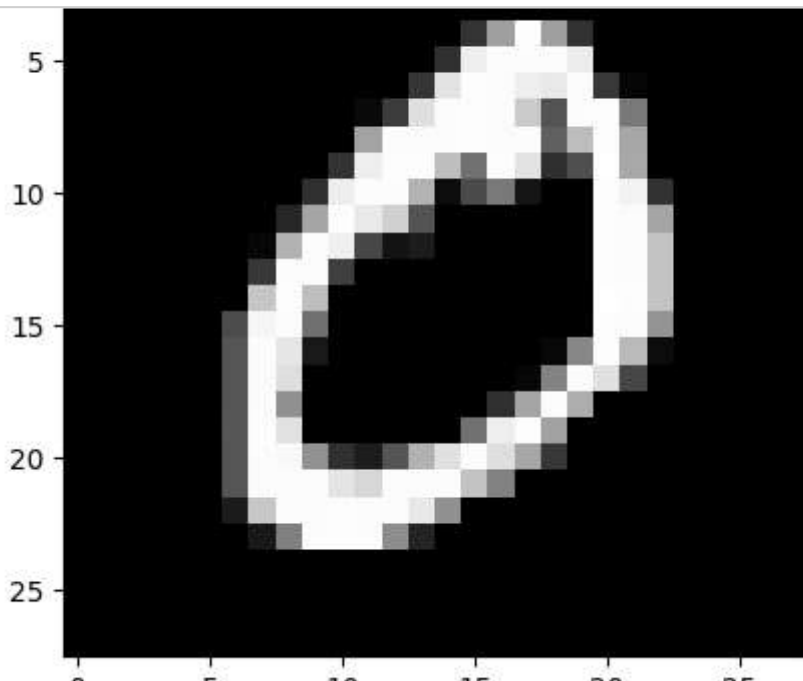
```
x = np.load("mnist.npy")
print(x.shape)
```

(70000, 28, 28)

Display the first few digits in the dataset.

In [16]:

```
for i in range(5):
    plt.imshow(x[i], cmap="gray")
    plt.show()
```



## Computing and diagonalizing the covariance of MNIST

We will interpret each image as a vector in  $\mathbb{R}^d$  with  $d = 28^2 = 768$ . The dataset can thus be seen as a matrix in  $\mathbb{R}^{n \times d}$  where  $n = 70000$ .

In [17]:

```
xx = x.reshape((x.shape[0], -1))
xx
```

Out[17]:

```
array([[0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       ...,
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0]], dtype=uint8)
```

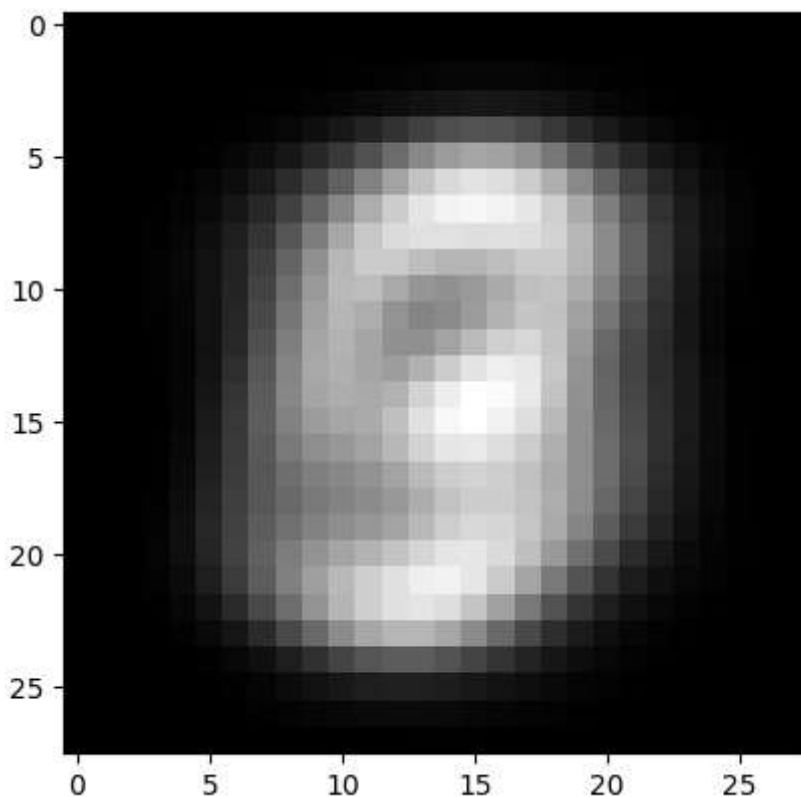
1. Compute the mean  $\mu \in \mathbb{R}^d$  of the MNIST dataset and plot it as a  $28 \times 28$  image.

In [18]:

```
# Your answer here
mean_vector = xx.mean(axis=0)
mean_image = mean_vector.reshape((28, 28))
plt.imshow(mean_image, cmap="gray")
```

Out[18]:

<matplotlib.image.AxesImage at 0x23ec5285fc0>



2. Compute the covariance  $\Sigma \in \mathbb{R}^{d \times d}$  of the MNIST dataset and diagonalize it using the function `np.linalg.eigh`.

In [19]:

# Your answer here

```
cov_matrix = (1 / xx.shape[0]) * (xx - mean_vector).T @ (xx - mean_vector)
eigenvalues, eigenvectors = np.linalg.eigh(cov_matrix)
eigenvalues, eigenvectors
```

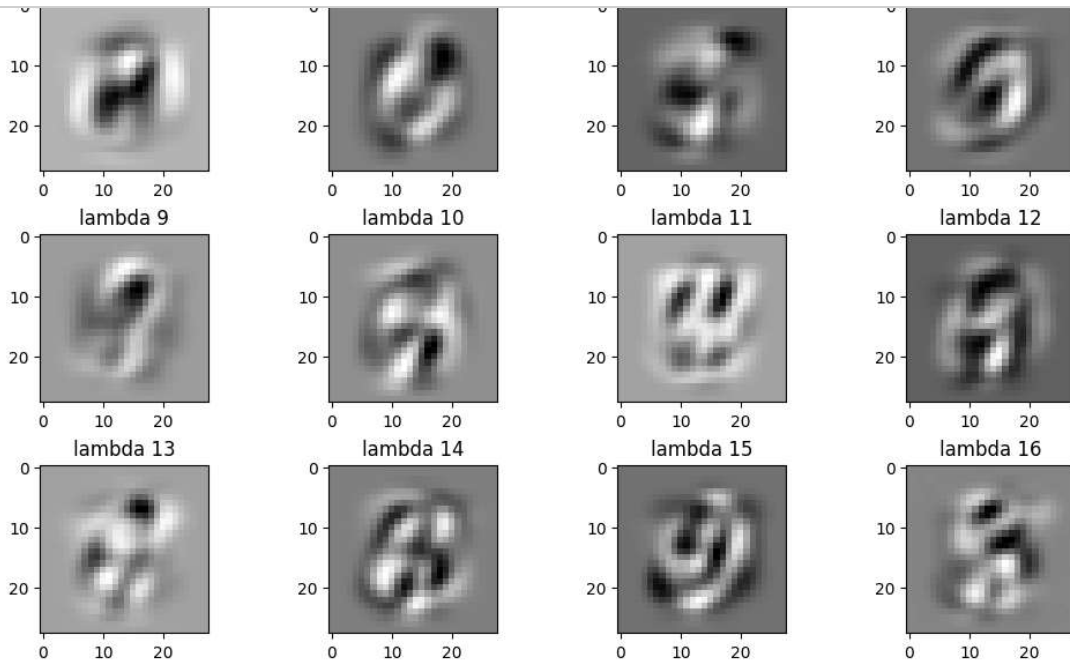
```
9.24113022e+03, 9.69015100e+03, 9.84750631e+03, 1.01692003e+04,
1.06378912e+04, 1.08680122e+04, 1.09654969e+04, 1.16184074e+04,
1.19714546e+04, 1.23968292e+04, 1.28893227e+04, 1.31610777e+04,
1.35883208e+04, 1.43447224e+04, 1.52605741e+04, 1.55847279e+04,
1.60383427e+04, 1.64280388e+04, 1.67066997e+04, 1.73116665e+04,
1.86409047e+04, 1.94395023e+04, 2.00863333e+04, 2.06079485e+04,
2.21394661e+04, 2.25055712e+04, 2.36673094e+04, 2.53908053e+04,
2.69499231e+04, 2.77770236e+04, 2.87712502e+04, 3.02965122e+04,
3.12002508e+04, 3.28986421e+04, 3.46356878e+04, 3.65648922e+04,
3.95454343e+04, 4.07231430e+04, 4.38698621e+04, 4.52536592e+04,
5.09812503e+04, 5.43096063e+04, 5.81044457e+04, 5.85518694e+04,
6.98875556e+04, 7.22588790e+04, 8.03348093e+04, 9.46111872e+04,
9.91139674e+04, 1.12443533e+05, 1.47668187e+05, 1.67689177e+05,
1.85334709e+05, 2.10927341e+05, 2.45429921e+05, 3.34289286e+05]),
array([[0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       ...,
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.]])
```

3. Plot the ordered eigenvalues  $\lambda_1 \geq \dots \geq \lambda_k \geq \dots$  as a function  $k = 1, \dots, d$  with the x axis in log scale, and the first few eigenvectors  $u_1, \dots, u_k, \dots$  as  $28 \times 28$  images.

In [21]:

```
# Your answer here
sorted_eigenvalues = sorted(eigenvalues, reverse=True)
sorted_eigenvectors = np.fliplr(eigenvectors)

plt.plot(range(1, len(sorted_eigenvalues)+1), sorted_eigenvalues)
fig, axes = plt.subplots(4, 4, layout='constrained', figsize=(10, 8))
for i in range(4):
    for j in range(4):
        index = 4 * i + j
        axes[i][j].imshow(sorted_eigenvectors[:, index].reshape(28, 28), cmap="gray")
        axes[i][j].set_title(f"lambda {index + 1}")
```



## PCA compression of MNIST

- Let  $k \in \mathbb{N}$ . Compute the  $k$ -dimensional PCA approximation  $z_1, \dots, z_n$  of the MNIST dataset using the eigenvectors  $u_1, \dots, u_k$ . Then, compute the reconstructed images  $\hat{x}_i = \mu + z_{i,1}u_1 + \dots + z_{i,k}u_k$ , which are equal to the mean  $\mu$  plus the orthogonal projection of  $x_i - \mu$  on  $\text{Span}(u_1, \dots, u_k)$ . Display the first 5 reconstructed images  $\hat{x}_1, \dots, \hat{x}_5$ . Choose a small value of  $k$  that still allows recognizing the digits.



In [10]:

```
# Your answer here
# Compute the dimension of each data point
def dim_data(data):
    return data.shape[1]

# Compute the mean of the sample
def mean_data(data):
    return np.mean(data, axis=0)

# Compute the standard deviation of the sample
def sd_data(data):
    return (data - np.mean(data, axis=0)) / np.std(data, axis=0)

# Centerize the data for further computation of the covariance matrix
def centerize_data(data):
    # Centerize the data
    return data - np.mean(data, axis=0)

# Compute the eigenbasis with k eigenvectors
def compute_eigenbasis_k(centered_data, k):
    # Compute the top k eigenvectors for our eigenbasis
    cov_matrix = 1 / len(centered_data) * centered_data.T @ centered_data

    eigenvalues, eigenvectors = np.linalg.eigh(cov_matrix)

    # Flip the columns, in reversed order.
    sorted_eigenvectors = np.fliplr(eigenvectors)[:,:k]

    return sorted_eigenvectors

# Main Procedure: PCA
def PCA_procedure(data, k):
    centered_data = centerize_data(data)
    V_k = compute_eigenbasis_k(centered_data, k)

    # Find PCA projection
    Z_k = V_k.T @ centered_data.T

    # Z_k is a matrix with (k, 70000), where the projection for each data point onto eigenbasis
    return Z_k, V_k

# Main Procedure: Inverse PCA
def inverse_PCA_procedure(Z_k, V_k, data, k):
    centered_data = centerize_data(data)

    mu = data.mean(axis=0)

    # Revert to origin
    RC_k = V_k @ Z_k + mu.reshape(dim_data(data), 1)

    # RC_k is a matrix with (784, 70000), where the reconstructed coordinates for each data point
    return RC_k
```

```

# Display the reconstructed images
def show_first_five_reconstructed(data, k):
    # Draw the first five reconstructed images
    fig, axes = plt.subplots(1, 5, figsize=(16,16), constrained_layout=True)

    Z_k, V_k = PCA_procedure(data, k)
    RC_k = inverse_PCA_procedure(Z_k, V_k, data, k)

    for i in range(5):
        axes[i].imshow(RC_k[:,i].reshape(28,28), cmap="gray")

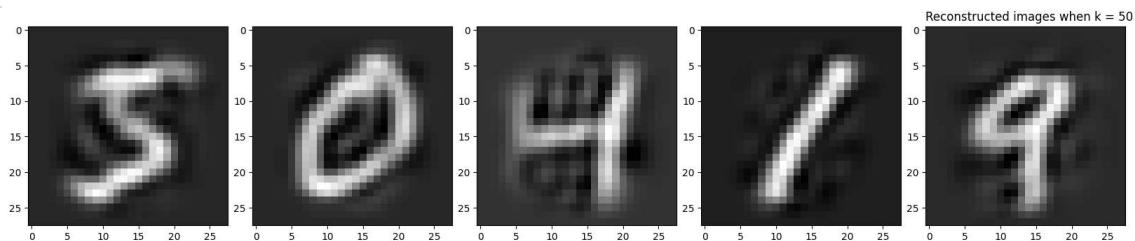
    plt.title(f"Reconstructed images when k = {k}", loc = "left")

def problem_1(data):
    k = 50
    show_first_five_reconstructed(data, k)

def problem_2(data):
    k_list = [10,30,50,100,300,500]
    for k in k_list:
        show_first_five_reconstructed(data, k)

# xx is (70000, 784)
problem_1(xx) # We pick 50, when we could recognize the reconstructed digits by raw eyes, which

```



In [11]:

```
problem_2(xx) # We find that 30 is minimum number of eigenvectors that are needed to reconstruct
```



In [ ]:



